

▼ Understanding Naïve Bayes Algorithm in Machine Learning

Naïve Bayes classification is a fundamental and widely used algorithm in machine learning, particularly in the domain of classification tasks.

The algorithm is based on applying Bayes' theorem with a key assumption of feature independence.

Bayes' Theorem:

Before delving into Naïve Bayes, let's understand Bayes' theorem, a fundamental principle in **probability theory** --> Multiplication rule -->

Independent event & Dependent event.

Bayes' theorem relates the conditional and marginal probabilities of events A and B, allowing us to calculate one given the other:

$$P(A/B) = P(B/A) * P(A) / P(B)$$

Where:

- **P(A/B)** is the probability of event A occurring given that event B has occurred.
- **P(B/A)** is the probability of event B occurring given that event A has occurred.
- **P(A)** is the prior probability of event A.
- **P(B)** is the prior probability of event B.

Application of Bayes' Theorem in Naïve Bayes Algorithm:

In the context of Naïve Bayes classification, we use Bayes' theorem to calculate the probability of a particular class given the features associated with a given instance. The algorithm assumes that the features are conditionally independent.

The formula for Naïve Bayes classification can be simplified using the assumption of feature independence:

$$P(y/X) = P(X/y) * P(y) / P(X)$$

Where:

- **P(y/X)** is the probability of class y given the features X (what we want to calculate).
- **P(X/y)** is the probability of observing features X given class y (likelihood).
- **P(y)** is the prior probability of class y.
- **P(X)** is the prior probability of features X (evidence).

Naïve Bayes Assumption of Feature Independence:

The "naïve" assumption in Naïve Bayes is that the features (attributes) used to describe the instances are independent of each other given the class label. This simplifies the calculation of the likelihood P(X/y) and makes the algorithm computationally efficient.

By assuming feature independence, we can represent the likelihood P(X/y) as the product of individual feature probabilities:

$$P(X/y) = P(y) * P(X1/y) * P(X2/y) * * P(Xn/y)$$

This simplification allows us to compute the nearer probability P(y/X) easily.

Why Use Naïve Bayes:

- Naïve Bayes is computationally efficient and easy to implement, making it suitable for large datasets.
- It's particularly effective for text-based data analysis, such as sentiment analysis, spam detection, document classification, etc.
- Naïve Bayes often performs well even with a small amount of training data.
- Despite its simplifying assumptions, Naïve Bayes can yield competitive classification performance compared to more complex algorithms.

▾ Types of Naive Bayes algorithm

There are 3 types of Naïve Bayes algorithm.

- Gaussian Naïve Bayes
- Multinomial Naïve Bayes
- Bernoulli Naïve Bayes

Gaussian Naïve Bayes Algorithm:

Gaussian Naïve Bayes designed for datasets with continuous attribute values. It makes the assumption that the values associated with each class are distributed according to a Gaussian (Normal) distribution.

we segment the training data by each class. Then compute the mean (μ) and variance (σ) of the continuous attribute values. Then, given a new observation (x_i), calculate the probability distribution of x_i for each class using the Gaussian distribution equation: https://scikit-learn.org/stable/modules/naive_bayes.html

This equation gives the probability of x_i belonging to a certain class based on its Gaussian distribution.

Multinomial Naïve Bayes Algorithm:

Multinomial Naïve Bayes is suitable for datasets where samples (feature vectors) represent the frequencies of events generated by a multinomial distribution. It is commonly used in text categorization, where features are typically word counts or term frequencies.

Bernoulli Naïve Bayes Algorithm:

Bernoulli Naïve Bayes is another variant suitable for datasets where features are binary or boolean variables, describing inputs as either present or absent. It's often used in document classification tasks.

These Naïve Bayes variants are widely used in various domains and provide efficient and effective solutions for classification tasks based on the underlying distribution of the data.

Applying Naïve Bayes Algorithm in Machine Learning

▾ Import Python Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

import seaborn as sns

Load Dataset

Dataset: <https://www.kaggle.com/datasets/qizarafzaal/adult-dataset/data>

Context of the Dataset : the dataset contains of adult population information and the thier income is more or less then 50k per year.

Goal: Using Gaussian Naïve Bayes to Predict whether income exceeds \$50K/yr based on census data.

```
df = pd.read_csv('/content/adult.csv', sep=',', header=None)
df.head()
```

```
<ipython-input-306-4c0a0d3d0859>:1: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as re
df = pd.read_csv('/content/adult.csv', sep=',', header=None)
```

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|---|----|------------------|--------|-----------|----|--------------------|-------------------|---------------|-------|--------|------|----|----|---------------|-------|
| 0 | 39 | State-gov | 77516 | Bachelors | 13 | Never-married | Adm-clerical | Not-in-family | White | Male | 2174 | 0 | 40 | United-States | <=50K |
| 1 | 50 | Self-emp-not-inc | 83311 | Bachelors | 13 | Married-civ-spouse | Exec-managerial | Husband | White | Male | 0 | 0 | 13 | United-States | <=50K |
| 2 | 38 | Private | 215646 | HS-grad | 9 | Divorced | Handlers-cleaners | Not-in-family | White | Male | 0 | 0 | 40 | United-States | <=50K |
| 3 | 53 | Private | 234721 | 11th | 7 | Married-civ-spouse | Handlers-cleaners | Husband | Black | Male | 0 | 0 | 40 | United-States | <=50K |
| 4 | 28 | Private | 338409 | Bachelors | 13 | Married-civ-spouse | Prof-specialty | Wife | Black | Female | 0 | 0 | 40 | Cuba | <=50K |

Set the Column name

```
column_names = ['age', 'workclass', 'fnlwt', 'education', 'education_num', 'marital_status', 'occupation', 'relationship',
                'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country', 'income']
```

```
df.columns = column_names
```

```
df.head()
```

| | age | workclass | fnlwt | education | education_num | marital_status | occupation | relationship | race | sex | capital_gain | capital_loss | hours_per_week | native_country | income |
|---|-----|------------------|--------|-----------|---------------|--------------------|-------------------|---------------|-------|--------|--------------|--------------|----------------|----------------|--------|
| 0 | 39 | State-gov | 77516 | Bachelors | 13 | Never-married | Adm-clerical | Not-in-family | White | Male | 2174 | 0 | 40 | United-States | <=50K |
| 1 | 50 | Self-emp-not-inc | 83311 | Bachelors | 13 | Married-civ-spouse | Exec-managerial | Husband | White | Male | 0 | 0 | 13 | United-States | <=50K |
| 2 | 38 | Private | 215646 | HS-grad | 9 | Divorced | Handlers-cleaners | Not-in-family | White | Male | 0 | 0 | 40 | United-States | <=50K |
| 3 | 53 | Private | 234721 | 11th | 7 | Married-civ-spouse | Handlers-cleaners | Husband | Black | Male | 0 | 0 | 40 | United-States | <=50K |
| 4 | 28 | Private | 338409 | Bachelors | 13 | Married-civ-spouse | Prof-specialty | Wife | Black | Female | 0 | 0 | 40 | Cuba | <=50K |

Exploratory Data Analysis

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt               32561 non-null  int64
3   education            32561 non-null  object
4   education_num        32561 non-null  int64
5   marital_status       32561 non-null  object
6   occupation           32561 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital_gain         32561 non-null  int64
11  capital_loss         32561 non-null  int64
12  hours_per_week       32561 non-null  int64
13  native_country       32561 non-null  object
14  income               32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

Types of Variable

```
catagorical_feature = df.select_dtypes(include='object').columns
numerical_feature = df.select_dtypes(exclude='object').columns
```

catagorical_feature

```
Index(['workclass', 'education', 'marital_status', 'occupation',
      'relationship', 'race', 'sex', 'native_country', 'income'],
      dtype='object')
```

numerical_feature

```
Index(['age', 'fnlwgt', 'education_num', 'capital_gain', 'capital_loss',
      'hours_per_week'],
      dtype='object')
```

df[catagorical_feature].isnull().sum()

```
workclass      0
education      0
marital_status  0
occupation     0
relationship   0
race           0
sex            0
native_country  0
income         0
dtype: int64
```

df[numerical_feature].isnull().sum()

```
age          0
fnlwgt       0
education_num 0
capital_gain 0
capital_loss 0
hours_per_week 0
dtype: int64
```

Exploring Catagorical Data

```
df['workclass'].unique()
```

```
array(['State-gov', 'Self-emp-not-inc', 'Private', 'Federal-gov',
       'Local-gov', '?', 'Self-emp-inc', 'Without-pay', 'Never-worked'],
      dtype=object)
```

```
df['workclass'].value_counts()
```

```
Private          22696
Self-emp-not-inc 2541
Local-gov        2093
?                1836
State-gov        1298
Self-emp-inc     1116
Federal-gov      960
Without-pay      14
Never-worked      7
Name: workclass, dtype: int64
```

```
df['education'].unique()
```

```
array(['Bachelors', 'HS-grad', '11th', 'Masters', '9th', 'Some-college',
       'Assoc-acdm', 'Assoc-voc', '7th-8th', 'Doctorate', 'Prof-school',
       '5th-6th', '10th', '1st-4th', 'Preschool', '12th'], dtype=object)
```

```
df['education'].value_counts()
```

```
HS-grad          10501
Some-college     7291
Bachelors        5355
Masters          1723
Assoc-voc        1382
11th             1175
Assoc-acdm       1067
10th             933
7th-8th          646
Prof-school      576
9th              514
12th             433
Doctorate        413
5th-6th          333
1st-4th          168
Preschool        51
Name: education, dtype: int64
```

```
df['marital_status'].unique()
```

```
array(['Never-married', 'Married-civ-spouse', 'Divorced',
       'Married-spouse-absent', 'Separated', 'Married-AF-spouse',
       'Widowed'], dtype=object)
```

```
df['marital_status'].value_counts()

Married-civ-spouse      14976
Never-married           10683
Divorced                 4443
Separated               1025
Widowed                 993
Married-spouse-absent   418
Married-AF-spouse       23
Name: marital_status, dtype: int64
```

```
df['occupation'].unique()
```

```
array(['Adm-clerical', 'Exec-managerial', 'Handlers-cleaners',
      'Prof-specialty', 'Other-service', 'Sales', 'Craft-repair',
      'Transport-moving', 'Farming-fishing', 'Machine-op-inspct',
      'Tech-support', '?', 'Protective-serv', 'Armed-Forces',
      'Priv-house-serv'], dtype=object)
```

```
df['occupation'].value_counts()
```

```
Prof-specialty      4140
Craft-repair        4099
Exec-managerial     4066
Adm-clerical        3770
Sales               3650
Other-service       3295
Machine-op-inspct   2002
?                  1843
Transport-moving    1597
Handlers-cleaners   1370
Farming-fishing     994
Tech-support        928
Protective-serv     649
Priv-house-serv     149
Armed-Forces        9
Name: occupation, dtype: int64
```

```
df['relationship'].unique()
```

```
array(['Not-in-family', 'Husband', 'Wife', 'Own-child', 'Unmarried',
      'Other-relative'], dtype=object)
```

```
df['relationship'].value_counts()
```

```
Husband      13193
Not-in-family 8305
Own-child    5068
Unmarried    3446
Wife         1568
Other-relative 981
Name: relationship, dtype: int64
```

```
df['race'].unique()
```

```
array(['White', 'Black', 'Asian-Pac-Islander', 'Amer-Indian-Eskimo',
      'Other'], dtype=object)
```

```
df['race'].value_counts()
```

```
White      27816
Black      3124
```

```
Asian-Pac-Islander      1039
Amer-Indian-Eskimo      311
Other                    271
Name: race, dtype: int64
```

```
df['sex'].unique()

array(['Male', 'Female'], dtype=object)
```

```
df['sex'].value_counts()

Male      21790
Female    10771
Name: sex, dtype: int64
```

```
df['native_country'].unique()

array(['United-States', 'Cuba', 'Jamaica', 'India', '?', 'Mexico',
       'South', 'Puerto-Rico', 'Honduras', 'England', 'Canada', 'Germany',
       'Iran', 'Philippines', 'Italy', 'Poland', 'Columbia', 'Cambodia',
       'Thailand', 'Ecuador', 'Laos', 'Taiwan', 'Haiti', 'Portugal',
       'Dominican-Republic', 'El-Salvador', 'France', 'Guatemala',
       'China', 'Japan', 'Yugoslavia', 'Peru',
       'Outlying-US(Guam-USVI-etc)', 'Scotland', 'Trinidad&Tobago',
       'Greece', 'Nicaragua', 'Vietnam', 'Hong', 'Ireland', 'Hungary',
       'Holand-Netherlands'], dtype=object)
```

```
df['native_country'].value_counts()

United-States      29170
Mexico              643
?                  583
Philippines        198
Germany            137
Canada             121
Puerto-Rico       114
El-Salvador        106
India              100
Cuba                95
England            90
Jamaica            81
South              80
China              75
Italy              73
Dominican-Republic 70
Vietnam            67
Guatemala          64
Japan              62
Poland             60
Columbia           59
Taiwan             51
Haiti              44
Iran               43
Portugal           37
Nicaragua          34
Peru               31
France             29
Greece             29
Ecuador            28
Ireland            24
Hong               20
Cambodia           19
Trinidad&Tobago    19
Laos               18
```

| | |
|----------------------------|----|
| Thailand | 18 |
| Yugoslavia | 16 |
| Outlying-US(Guam-USVI-etc) | 14 |
| Honduras | 13 |
| Hungary | 13 |
| Scotland | 12 |
| Holand-Netherlands | 1 |

Name: native_country, dtype: int64

```
df['income'].unique()

array(['<=50K', '>50K'], dtype=object)
```

```
df['income'].value_counts()

<=50K    24720
>50K      7841
Name: income, dtype: int64
```

From the above exploration into Catagorical data, we can see that variables **workclass**, **occupation** and **native_country** contain (?) which is missing values.

To detect the missing values generally we see it contains NaN after python code `df.isnull().sum()`, but we did not see because ? does not cosider as missing in python.

So, we replace ? into NaN.

Repalcing '?' into NaN

```
df['workclass'].replace('?', np.NaN, inplace=True)
```

```
df['occupation'].replace('?', np.NaN, inplace=True)
```

```
df['native_country'].replace('?', np.NaN, inplace=True)
```

```
df['occupation'].value_counts()

Prof-specialty    4140
Craft-repair      4099
Exec-managerial   4066
Adm-clerical      3770
Sales             3650
Other-service     3295
Machine-op-inspct 2002
Transport-moving  1597
Handlers-cleaners 1370
Farming-fishing   994
Tech-support      928
Protective-serv   649
Priv-house-serv   149
Armed-Forces      9
Name: occupation, dtype: int64
```

```
df[catagorical_feature].isnull().sum()

workclass    1836
education      0
marital_status  0
```



```
occupation      1843
relationship     0
race            0
sex            0
native_country  583
income          0
dtype: int64
```

Cardinality

Low cardinality: If there's only one category in a column, it won't provide any unique information to our model. Low cardinality means the observation in the columns has constant value which means same value in all the rows of the columns. Ex. type of building has only apartment. So don't include in the model. Drop

High cardinality: as low cardinality don't give you same information to the model, High cardinality doesn't give any information to the model. Can drop also.

As low cardinality gives low or no information and high cardinality is overload with information which both doesn't help the model since model looks for trend.

```
for col in catagorical_feature:
    print(col, len(df[col].unique()))

workclass 9
education 16
marital_status 7
occupation 15
relationship 6
race 5
sex 2
native_country 42
income 2
```

Numerical feature

```
df[numerical_feature].describe().T
```

| | count | mean | std | min | 25% | 50% | 75% | max |
|----------------|---------|---------------|---------------|---------|----------|----------|----------|-----------|
| age | 32561.0 | 38.581647 | 13.640433 | 17.0 | 28.0 | 37.0 | 48.0 | 90.0 |
| fnlwgt | 32561.0 | 189778.366512 | 105549.977697 | 12285.0 | 117827.0 | 178356.0 | 237051.0 | 1484705.0 |
| education_num | 32561.0 | 10.080679 | 2.572720 | 1.0 | 9.0 | 10.0 | 12.0 | 16.0 |
| capital_gain | 32561.0 | 1077.648844 | 7385.292085 | 0.0 | 0.0 | 0.0 | 0.0 | 99999.0 |
| capital_loss | 32561.0 | 87.303830 | 402.960219 | 0.0 | 0.0 | 0.0 | 0.0 | 4356.0 |
| hours_per_week | 32561.0 | 40.437456 | 12.347429 | 1.0 | 40.0 | 40.0 | 45.0 | 99.0 |

Split

A key part in any model-building project is separating your target (y) (the thing you want to predict) from your features (X) (the information your model will use to make its predictions).

```
X = df.drop(['income'], axis=1)
y = df['income']
```

Train Test split

```
# split X and y into training and testing sets
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
```

```
print('X_train:',X_train.shape)
print('y_train:',y_train.shape)
print('X_test:',X_test.shape)
print('y_test:',y_test.shape)
```

```
    X_train: (22792, 14)
    y_train: (22792,)
    X_test: (9769, 14)
    y_test: (9769,)
```

▼ Feature Engineering

Feature Engineering is the process of transforming raw data into useful features that help us to understand our model better and increase its predictive power. I will carry out feature engineering on different types of variables.

- Imputaion
- Encoding
- Scaling

```
cat_col = X_train.select_dtypes(include='object').columns
num_col = X_train.select_dtypes(exclude='object').columns
```

Handling the Missing values

```
X_train[num_col].isnull().sum()
```

```
    age          0
    fnlwgt        0
    education_num  0
    capital_gain  0
    capital_loss  0
    hours_per_week 0
    dtype: int64
```

```
X_train[cat_col].isnull().sum()
```

```
    workclass      1276
    education        0
    marital_status  0
    occupation     1278
    relationship    0
    race            0
    sex             0
    native_country  414
    dtype: int64
```

Imputatiuon

There are two methods can be used to impute missing values.

- mean or median or mode imputation
- random sample imputation

When there are outliers in the dataset, we should use median imputation.

impute missing categorical variables with **most frequent value**

SimpleImputer is a scikit-learn class which is helpful in handling the missing data

implemented by the use of the

- SimpleImputer():
 - missing_values : The missing_values placeholder which has to be imputed. By default is NaN
 - strategy: The data which will replace the NaN values from the dataset. The strategy argument can take the values – ‘mean’(default), ‘median’, ‘most_frequent’ and ‘constant’.

```
from sklearn.impute import SimpleImputer

# Define columns with missing values
columns_with_missing_values = ['workclass', 'occupation', 'native_country']

# Create the SimpleImputer object
imputer = SimpleImputer(strategy='most_frequent')

# Fit and transform X_train
X_train[columns_with_missing_values] = imputer.fit_transform(X_train[columns_with_missing_values])

# Transform X_test
X_test[columns_with_missing_values] = imputer.transform(X_test[columns_with_missing_values])
```

```
X_test[cat_col].isnull().sum()
```

```
workclass      0
education      0
marital_status  0
occupation     0
relationship    0
race           0
sex            0
native_country  0
dtype: int64
```

```
X_train[cat_col].isnull().sum()
```

```
workclass      0
education      0
marital_status  0
occupation     0
relationship    0
race           0
sex            0
native_country  0
dtype: int64
```

X_train.head()

| | age | workclass | fnlwgt | education | education_num | marital_status | occupation | relationship | race | sex | capital_gain | capital_loss | hours_per_week | native_country |
|-------|-----|-----------|--------|--------------|---------------|--------------------|--------------|---------------|-------|--------|--------------|--------------|----------------|----------------|
| 32098 | 45 | Private | 170871 | HS-grad | 9 | Married-civ-spouse | Craft-repair | Husband | White | Male | 7298 | 0 | 60 | United-States |
| 25206 | 47 | State-gov | 108890 | HS-grad | 9 | Divorced | Adm-clerical | Unmarried | White | Female | 1831 | 0 | 38 | United-States |
| 23491 | 48 | Private | 187505 | Some-college | 10 | Married-civ-spouse | Sales | Husband | White | Male | 0 | 0 | 50 | United-States |
| 12367 | 29 | Private | 145592 | HS-grad | 9 | Never-married | Craft-repair | Not-in-family | White | Male | 0 | 0 | 40 | Guatemala |
| 7054 | 23 | Private | 203003 | 7th-8th | 4 | Never-married | Craft-repair | Not-in-family | White | Male | 0 | 0 | 25 | Germany |

Encoding

OHE is the standard approach to encode categorical data.

One hot encoding (OHE) creates a binary variable for each one of the different categories present in a variable. These binary variables take 1 if the observation shows a certain category or 0 otherwise. OHE is suitable for linear models.

One hot encoding (OHE) creates by replacing the categorical variable by different boolean variables, which take value 0 or 1, to indicate whether or not a certain category / label of the variable was present for that observation. Each one of the boolean variables are also known as dummy variables or binary variables.

For example, from the categorical variable "Gender", with labels 'female' and 'male', we can generate the boolean variable "female", which takes 1 if the person is female or 0 otherwise. We can also generate the variable male, which takes 1 if the person is "male" and 0 otherwise.

```
import category_encoders as ce

encoder = ce.OneHotEncoder(cols=['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country'])

X_train = encoder.fit_transform(X_train)

X_test = encoder.transform(X_test)
```

```
from sklearn.preprocessing import OneHotEncoder
# Assuming 'cat_col' contains the names of categorical columns
encoder = OneHotEncoder()

# Fit and transform X_train
X_train_encoded = encoder.fit_transform(X_train[cat_col])
X_train_encoded_df = pd.DataFrame(X_train_encoded.toarray(), columns=encoder.get_feature_names_out(cat_col))

# Transform X_test
X_test_encoded = encoder.transform(X_test[cat_col])
X_test_encoded_df = pd.DataFrame(X_test_encoded.toarray(), columns=encoder.get_feature_names_out(cat_col))

# Drop the original categorical columns
X_train = X_train.drop(cat_col, axis=1)
X_test = X_test.drop(cat_col, axis=1)

# Concatenate the encoded columns
X_train = pd.concat([X_train.reset_index(drop=True), X_train_encoded_df], axis=1)
X_test = pd.concat([X_test.reset_index(drop=True), X_test_encoded_df], axis=1)
```

```
X_test = pd.concat([X_test.reset_index(drop=True), X_test_encoded_df], axis=1)
```

```
X_train.shape

(22792, 105)

X_test.shape

(9769, 105)
```

Feature Scaling

- StandardScaler
- MinMaxScaler
- RobustScaler

RobustScaler is a method for scaling features in a dataset using statistics that are robust to outliers.

When you scale features using RobustScaler, it removes the median and scales the data based on the interquartile range (IQR). This scaling is more robust to outliers compared to standard scaling methods like mean and variance scaling.

```
cols = X_train.columns

from sklearn.preprocessing import RobustScaler
scaler = RobustScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

X_train = pd.DataFrame(X_train, columns=[cols])
X_test = pd.DataFrame(X_test, columns=[cols])
```

```
X_train.head()
```

| | age | fnlwgt | education_num | capital_gain | capital_loss | hours_per_week | workclass_Federal- gov | workclass_Local- gov | workclass_Never- worked | workclass_Private | ... | native_country_Portugal | native_country_Puerto- Rico | native_country_Other |
|---|-------|-----------|---------------|--------------|--------------|----------------|---------------------------|-------------------------|----------------------------|-------------------|-----|-------------------------|--------------------------------|----------------------|
| 0 | 0.40 | -0.058906 | -0.333333 | 7298.0 | 0.0 | 4.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 |
| 1 | 0.50 | -0.578076 | -0.333333 | 1831.0 | 0.0 | -0.4 | 0.0 | 0.0 | 0.0 | -1.0 | ... | 0.0 | 0.0 | 0.0 |
| 2 | 0.55 | 0.080425 | 0.000000 | 0.0 | 0.0 | 2.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 |
| 3 | -0.40 | -0.270650 | -0.333333 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 |
| 4 | -0.70 | 0.210240 | -2.000000 | 0.0 | 0.0 | -3.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 |

5 rows × 105 columns

Building Model

Baseline: The first step in building a model is baselining. To do this, ask yourself how you will know if the model you build is performing well?

Model Training

The steps to building and using a model are:

Define: What type of model will it be? A decision tree? Some other type of model? Some other parameters of the model type are specified too.

Fit: Capture patterns from provided data. This is the heart of modeling.

Predict: Just what it sounds like

Evaluate: Determine how accurate the model's predictions are.

```
# train a Gaussian Naive Bayes classifier on the training set
from sklearn.naive_bayes import GaussianNB
```

```
# instantiate the model
gnb = GaussianNB()
```

```
# fit the model
gnb.fit(X_train, y_train)
```

```
▼ GaussianNB
GaussianNB()
```

Predict the Model

```
y_pred = gnb.predict(X_test)
```

```
y_pred
```

```
array(['<=50K', '<=50K', '>50K', ..., '>50K', '<=50K', '<=50K'],
      dtype='<U5')
```

Predicting test set result

At this point, the model is now trained and ready to predict the output of new observations. Remember, we split our dataset into train and test sets. We will provide test sets to the model and check its performance.

#y_test and y_pred are your actual and predicted labels

```
prediction_df = pd.DataFrame({
    'Actual Value': y_test,
    'Predicted Value': y_pred,
    'Prediction Correct': y_test == y_pred # True if prediction is correct, False otherwise
})
```

```
# Display the prediction_df DataFrame
print(prediction_df)
```

| | Actual Value | Predicted Value | Prediction Correct |
|-------|--------------|-----------------|--------------------|
| 22278 | <=50K | <=50K | True |
| 8950 | <=50K | <=50K | True |
| 7838 | <=50K | >50K | False |
| 16505 | <=50K | >50K | False |
| 19140 | >50K | >50K | True |
| ... | ... | ... | ... |
| 21949 | >50K | >50K | True |
| 26405 | >50K | >50K | True |
| 23236 | >50K | >50K | True |
| 26823 | <=50K | <=50K | True |
| 20721 | <=50K | <=50K | True |

[9769 rows x 3 columns]

▼ Evaluate the Model

accuracy score

```
from sklearn.metrics import accuracy_score
print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred)))
```

```
Model accuracy score: 0.8083
```

Null or Baseline accuracy

Comparing model accuracy with a null or baseline accuracy is a good practice to evaluate the model's performance. The null accuracy is the accuracy achieved by a model that always predicts the most frequent class in the dataset. It provides a baseline for comparison, helping to gauge whether the model's performance is meaningful and better than a simple baseline prediction strategy.

```
y_test.value_counts()
```

```
<=50K    7407
>50K     2362
Name: income, dtype: int64
```

```
baseline_accuracy = (7407/(7407+2362))
```

```
print('Baseline accuracy score: {0:0.4f}'.format(baseline_accuracy))
```

```
Baseline accuracy score: 0.7582
```

Check for overfitting and underfitting:

- Overfitting usually manifests as a significant gap between training and test accuracies.
- Underfitting is marked by low accuracies on both sets due to insufficient model complexity.
- Generalized Model is demonstrates consistent performance on both training and test sets, suggesting it is well-generalized and not overfit.

```
print('Training set score: {:.4f}'.format(gnb.score(X_train, y_train)))
```

```
print('Test set score: {:.4f}'.format(gnb.score(X_test, y_test)))
```

```
Training set score: 0.8067
Test set score: 0.8083
```

training set= 0.8067 and test set=0.8083 , it means that the model is performing consistently well on both the training and test data. The scores are close to each other, indicating that the model is likely well-generalized and not overfitting.

Based on the model classification accuracy and baseline accuracy, and generalized model, we can say model is performing very good. Our model is able to predict the class labels.

But, it does not give the underlying distribution of values and it does not tell anything about the type of errors our classifier is making. For that we use Confusion Matrix

Confusion Matrix A confusion matrix is a tool for summarizing the performance of a classification algorithm. A confusion matrix will give us a clear picture of classification model performance and the types of errors produced by the model. It gives us a summary of correct and incorrect predictions broken down by each category. The summary is represented in a tabular form.

Four types of outcomes are possible while evaluating a classification model performance. These four outcomes are described below:-

- 1. True Positives (TP) – True Positives occur when we predict an observation belongs to a certain class and the observation actually belongs to that class.
- 2. True Negatives (TN) – True Negatives occur when we predict an observation does not belong to a certain class and the observation actually does not belong to that class.
- 3. False Positives (FP) – False Positives occur when we predict an observation belongs to a certain class but the observation actually does not belong to that class. This type of error is called Type I error.
- 4. False Negatives (FN) – False Negatives occur when we predict an observation does not belong to a certain class but the observation actually belongs to that class. This is a very serious error and it is called Type II error.

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
```

```
print('Confusion matrix\n\n', cm)
```

```
print('\nTrue Positives(TP) = ', cm[0,0])
print('\nTrue Negatives(TN) = ', cm[1,1])
print('\nFalse Positives(FP) = ', cm[0,1])
print('\nFalse Negatives(FN) = ', cm[1,0])
```

Confusion matrix

| | | |
|-----------------------|--------------|--|
| | | |
| | [[5999 1408] | |
| | [465 1897]] | |
| True Positives(TP) = | 5999 | |
| True Negatives(TN) = | 1897 | |
| False Positives(FP) = | 1408 | |
| False Negatives(FN) = | 465 | |

The confusion matrix shows 5999 + 1897 = 7896 correct predictions
and 1408 + 465 = 1873 incorrect predictions.

In this case, we have

True Positives (Actual Positive:1 and Predict Positive:1) - 5999

True Negatives (Actual Negative:0 and Predict Negative:0) - 1897

False Positives (Actual Negative:0 but Predict Positive:1) - 1408 (Type I error)

False Negatives (Actual Positive:1 but Predict Negative:0) - 465 (Type II error)

Classification metrices

Classification Report

Classification report is another way to evaluate the classification model performance. It displays the precision, recall, f1 and support scores for the model.

```
from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| <=50K | 0.93 | 0.81 | 0.86 | 7407 |
| >50K | 0.57 | 0.80 | 0.67 | 2362 |
| accuracy | | | 0.81 | 9769 |
| macro avg | 0.75 | 0.81 | 0.77 | 9769 |
| weighted avg | 0.84 | 0.81 | 0.82 | 9769 |

Precision

Precision can be defined as the percentage of correctly predicted positive outcomes out of all the predicted positive outcomes. It can be given as the ratio of true positives (TP) to the sum of true and false positives (TP + FP).

precision = TP / float(TP + FP)

Recall

Recall can be defined as the percentage of correctly predicted positive outcomes out of all the actual positive outcomes. It can be given as the ratio of true positives (TP) to the sum of true positives and false negatives (TP + FN). Recall is also called Sensitivity.

recall = TP / float(TP + FN)

f1-score

f1-score is the weighted harmonic mean of precision and recall. The best possible f1-score would be 1.0 and the worst would be 0.0.

f1-score is the harmonic mean of precision and recall. So, f1-score is always lower than accuracy measures as they embed precision and recall into their computation. The weighted average of f1-score should be used to compare classifier models, not global accuracy.

Support

Support is the actual number of occurrences of the class in our dataset.

▼ **k-Fold Cross Validation**

```
# Applying 10-Fold Cross Validation

from sklearn.model_selection import cross_val_score

scores = cross_val_score(gnb, X_train, y_train, cv = 10, scoring='accuracy')

print('Cross-validation scores:{}'.format(scores))

Cross-validation scores:[0.81359649 0.80438596 0.81175954 0.8056165  0.79596314 0.79684072
0.81044318 0.81175954 0.80210619 0.81044318]
```

We can summarize the cross-validation accuracy by calculating its mean.

```
print('Average cross-validation score: {:.4f}'.format(scores.mean()))
```

```
Average cross-validation score: 0.8063
```

original model accuracy is 0.8083, but the mean cross-validation accuracy is 0.8063.